

lab-2-series-de-tiempo

August 10, 2023

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.tsa as tsa
import statsmodels as sm
from datetime import datetime
from statsmodels.tsa.stattools import adfuller
from statsmodels.tsa.seasonal import seasonal_decompose
from statsmodels.tsa.stattools import acf, pacf
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.tsa.statespace.sarimax import SARIMAX

from sklearn.preprocessing import PowerTransformer, StandardScaler
from sklearn.metrics import mean_squared_error, mean_absolute_error
from keras.models import Sequential
from keras.layers import LSTM, Dense
```

```
[2]: df_imp = pd.read_excel('IMPORTACION-VOLUMEN-2023-05.xlsx',
    sheet_name='IMPORTACION', skiprows = 6,
    nrows=270, header=0, usecols= 'A, G, I, K, L, U',
    engine='openpyxl').dropna()

df_imp.head()
```

```
[2]:
```

	Fecha	Diesel bajo azufre	Gas licuado de petróleo	Gasolina regular \
0	2001-01-01	0.0	194065.738095	177776.50
1	2001-02-01	0.0	170703.380952	123115.99
2	2001-03-01	0.0	161837.371429	161726.42
3	2001-04-01	0.0	163048.642857	127338.74
4	2001-05-01	0.0	171518.861905	168730.19

	Gasolina superior	Diesel alto azufre
0	373963.96	566101.99
1	243091.07	489525.80
2	312084.38	575559.68

3	285054.89	437745.42
4	300913.67	552609.13

```
[3]: df_imp.tail()
```

```
[3]:
```

	Fecha	Diesel bajo azufre	Gas licuado de petróleo	Gasolina regular \
264	2023-01-01	1442099.08	497780.69	909391.13
265	2023-02-01	1267967.39	652984.71	725101.20
266	2023-03-01	1317519.91	711978.35	803262.67
267	2023-04-01	1417182.73	647666.30	922032.39
268	2023-05-01	1428099.62	713348.99	947633.29

	Gasolina superior	Diesel alto azufre
264	578792.14	0.0
265	685183.06	0.0
266	633849.05	0.0
267	572201.36	0.0
268	668478.73	0.0

```
[4]: (df_imp.iloc[:,1] != 0).idxmax(), (df_imp.iloc[:,5] == 0).idxmax() # las
      ↪ columnas se pueden sumar en una sola
```

```
[4]: (204, 204)
```

```
[5]: df_imp['Diesel'] = df_imp['Diesel bajo azufre'] + df_imp['Diesel alto azufre']
      df_imp.drop(['Diesel bajo azufre', 'Diesel alto azufre'], axis=1, inplace=True)
```

```
[6]: df_imp.info(), df_imp.shape
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 269 entries, 0 to 268
Data columns (total 5 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Fecha                                269 non-null    datetime64[ns]
1   Gas licuado de petróleo              269 non-null    float64
2   Gasolina regular                    269 non-null    float64
3   Gasolina superior                   269 non-null    float64
4   Diesel                              269 non-null    float64
dtypes: datetime64[ns](1), float64(4)
memory usage: 12.6 KB
```

```
[6]: (None, (269, 5))
```

```
[7]: df_imp.describe().applymap(lambda x: f"{int(x):,}")
```

```
[7]:
```

	Gas licuado de petróleo	Gasolina regular	Gasolina superior	Diesel
count	269	269	269	269
mean	399,026	381,883	478,288	863,979
std	190,173	219,364	153,277	276,545
min	100,561	81,015	170,292	229,764
25%	218,257	201,907	363,352	678,749
50%	396,363	305,285	456,588	824,047
75%	540,671	528,347	582,374	1,050,560
max	960,840	987,872	1,227,173	1,595,698

```
[8]: df_p_nac_1 = pd.read_excel('Precios-Promedio-Nacionales-Diarios-2023.xlsx',
    ↪sheet_name='2021',skiprows = 6,
    ↪nrows=366, header=0,usecols= 'A,C:E, G',
    ↪engine='openpyxl')
df_p_nac_1.drop(0, inplace=True)
df_p_nac_2 = pd.read_excel('Precios-Promedio-Nacionales-Diarios-2023.xlsx',
    ↪sheet_name='2022',skiprows = 6,
    ↪nrows=366, header=0,usecols= 'A,C:E, G',
    ↪engine='openpyxl')
df_p_nac_2.drop(0, inplace=True)
df_p_nac_3 = pd.read_excel('Precios-Promedio-Nacionales-Diarios-2023.xlsx',
    ↪sheet_name='2023',skiprows = 7,
    ↪nrows=213, header=0,usecols= 'A,C:E, G',
    ↪engine='openpyxl')
df_p_nac_3.drop(0, inplace=True)
```

```
[9]: df_p_nac_1.head()
```

```
[9]:
```

	FECHA	Superior	Regular	Diesel	Glp Cilindro 25Lbs.
1	2021-01-01	21.91	21.11	17.61	99
2	2021-01-02	21.91	21.11	17.61	99
3	2021-01-03	21.91	21.11	17.61	99
4	2021-01-04	21.91	21.11	17.61	99
5	2021-01-05	21.91	21.11	17.61	99

```
[10]: df_p_nac_1.tail()
```

```
[10]:
```

	FECHA	Superior	Regular	Diesel	Glp Cilindro 25Lbs.
361	2021-12-27	28.69	27.91	24.51	122
362	2021-12-28	28.69	27.91	24.51	122
363	2021-12-29	28.69	27.91	24.51	122
364	2021-12-30	28.69	27.91	24.51	122
365	2021-12-31	28.69	27.91	24.51	122

```
[11]: df_p_nac_2.head()
```

```
[11]:          FECHA Superior Regular Diesel Glp Cilindro 25Lbs.
1 2022-01-01    28.69    27.91    24.51          122
2 2022-01-02    28.69    27.91    24.51          122
3 2022-01-03    28.79    27.99    24.6          122
4 2022-01-04    28.79    27.99    24.6          122
5 2022-01-05    28.79    27.99    24.6          122
```

```
[12]: df_p_nac_2.tail()
```

```
[12]:          FECHA Superior Regular Diesel Glp Cilindro 25Lbs.
361 2022-12-27    32.55    30.99    33.75          122
362 2022-12-28    32.55    30.99    33.75          122
363 2022-12-29    32.55    30.99    33.75          122
364 2022-12-30    32.55    30.99    33.75          122
365 2022-12-31    32.55    30.99    33.75          122
```

```
[13]: df_p_nac_3.head()
```

```
[13]:          FECHA Superior Regular Diesel Glp Cilindro 25Lbs.
1 2023-01-01    32.55    30.99    33.75          122
2 2023-01-02    32.55    30.99    33.75          122
3 2023-01-03    32.56    31.42    35.31          122
4 2023-01-04    32.56    31.42    35.31          122
5 2023-01-05    32.56    31.42    35.31          122
```

```
[14]: df_p_nac_3.tail()
```

```
[14]:          FECHA Superior Regular Diesel Glp Cilindro 25Lbs.
208 2023-07-27    34.78    33.28    28.47          135
209 2023-07-28    34.78    33.28    28.47          135
210 2023-07-29    34.78    33.28    28.47          135
211 2023-07-30    34.78    33.28    28.47          135
212 2023-07-31    36.29    34.77    29.68          135
```

```
[15]: df_precios = pd.concat([df_p_nac_1, df_p_nac_2, df_p_nac_3], ignore_index=True)
df_precios.head()
```

```
[15]:          FECHA Superior Regular Diesel Glp Cilindro 25Lbs.
0 2021-01-01    21.91    21.11    17.61          99
1 2021-01-02    21.91    21.11    17.61          99
2 2021-01-03    21.91    21.11    17.61          99
3 2021-01-04    21.91    21.11    17.61          99
4 2021-01-05    21.91    21.11    17.61          99
```

```
[16]: df_precios.tail()
```

```
[16]:          FECHA Superior Regular Diesel Glp Cilindro 25Lbs.
937 2023-07-27    34.78    33.28    28.47                135
938 2023-07-28    34.78    33.28    28.47                135
939 2023-07-29    34.78    33.28    28.47                135
940 2023-07-30    34.78    33.28    28.47                135
941 2023-07-31    36.29    34.77    29.68                135
```

```
[17]: df_precios.info(), df_precios.shape
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 942 entries, 0 to 941
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   FECHA                 942 non-null   datetime64[ns]
1   Superior              942 non-null   object
2   Regular               942 non-null   object
3   Diesel                942 non-null   object
4   Glp Cilindro 25Lbs.   942 non-null   object
dtypes: datetime64[ns](1), object(4)
memory usage: 36.9+ KB
```

```
[17]: (None, (942, 5))
```

```
[18]: columnas_precios = list(df_precios.columns[1:])
columnas_precios
```

```
[18]: ['Superior', 'Regular', 'Diesel', 'Glp Cilindro 25Lbs.']
```

```
[19]: df_precios[columnas_precios] = df_precios[columnas_precios].astype(float)
df_precios.info(), df_precios.shape
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 942 entries, 0 to 941
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   FECHA                 942 non-null   datetime64[ns]
1   Superior              942 non-null   float64
2   Regular               942 non-null   float64
3   Diesel                942 non-null   float64
4   Glp Cilindro 25Lbs.   942 non-null   float64
dtypes: datetime64[ns](1), float64(4)
memory usage: 36.9 KB
```

```
[19]: (None, (942, 5))
```

```
[20]: df_precios.describe()
```

```
[20]:
```

	Superior	Regular	Diesel	Glp Cilindro 25Lbs.
count	942.000000	942.000000	942.000000	942.000000
mean	32.339915	31.185817	28.495563	123.651921
std	4.555857	4.225467	6.059121	8.420853
min	21.910000	21.110000	17.610000	99.000000
25%	28.790000	27.990000	23.090000	120.000000
50%	33.100000	31.800000	27.680000	122.000000
75%	35.550000	34.230000	33.847500	122.000000
max	43.240000	40.500000	41.270000	147.000000

```
[21]: df_precios.iloc[:,1:].describe().applymap(lambda x: f"{int(x):,}")
```

```
[21]:
```

	Superior	Regular	Diesel	Glp Cilindro 25Lbs.
count	942	942	942	942
mean	32	31	28	123
std	4	4	6	8
min	21	21	17	99
25%	28	27	23	120
50%	33	31	27	122
75%	35	34	33	122
max	43	40	41	147

```
[22]: df_consumo = pd.read_excel('CONSUMO-2023-05.xlsx',
    ↳ sheet_name='CONSUMO', skiprows = 5,
    ↳ nrows=282, header=1, usecols= 'A, G, K, L, U',
    ↳ engine='openpyxl').dropna()
df_consumo.tail()
```

```
[22]:
```

	Fecha	Diesel bajo azufre	Gasolina regular	Gasolina superior \
276	2023-01-01	1285932.51	804524.73	599501.29
277	2023-02-01	1226884.63	737345.64	544144.00
278	2023-03-01	1463008.43	875010.34	656941.64
279	2023-04-01	1248429.37	798128.36	585642.20
280	2023-05-01	1346554.12	866826.79	646221.20

	Diesel alto azufre
276	0.0
277	0.0
278	0.0
279	0.0
280	0.0

```
[23]: (df_consumo.iloc[:,1] != 0).idxmax(), (df_consumo.iloc[:,4] == 0).idxmax()
```

```
[23]: (216, 216)
```

```
[24]: df_consumo['Diesel'] = df_consumo['Diesel bajo azufre'] + df_consumo['Diesel_
      ↪alto azufre']
      df_consumo.drop(['Diesel bajo azufre', 'Diesel alto azufre'], axis=1,
      ↪inplace=True)
```

```
[25]: df_consumo.head()
```

```
[25]:
```

	Fecha	Gasolina regular	Gasolina superior	Diesel
0	2000-01-01	202645.20	308156.82	634667.06
1	2000-02-01	205530.96	307766.31	642380.66
2	2000-03-01	229499.56	331910.29	699807.25
3	2000-04-01	210680.40	315648.08	586803.98
4	2000-05-01	208164.34	319667.97	656948.20


```
[26]: df_imp.set_index('Fecha', inplace=True)
      df_precios.set_index('FECHA', inplace=True)
      df_consumo.set_index('Fecha', inplace=True)
```

0.1 Descripción de las series de tiempo

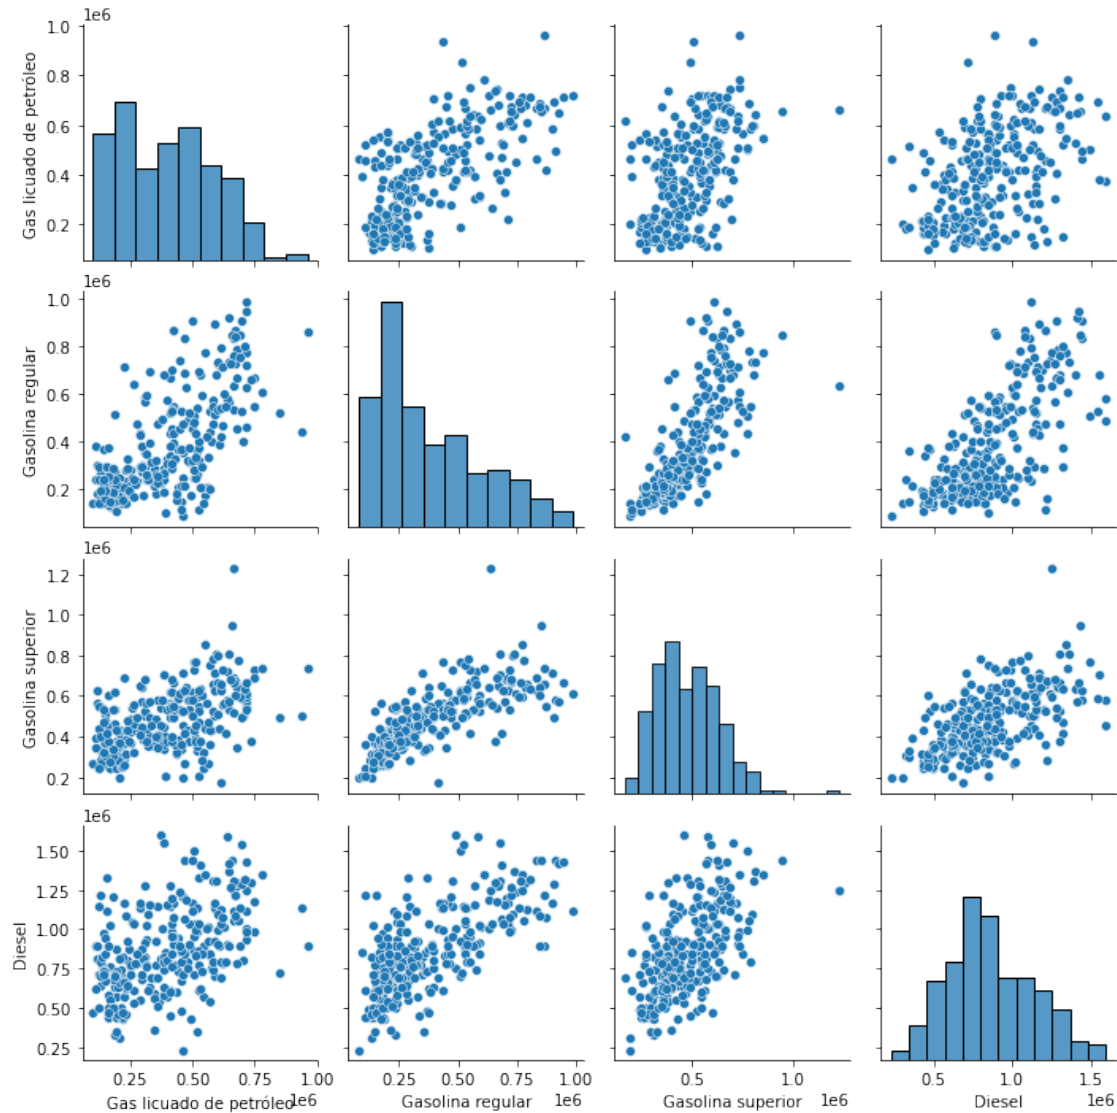
Se importaron tres archivos: 1. De “IMPORTACION-VOLUMEN-2023-05” se importaron los volúmenes de importación de combustibles de las columnas Diesel Bajo en Azufre, Gas Licuado de Petróleo, Gasolina Regular, Gasolina Superior y Diesel alto en Azufre calculado de forma mensual de Enero 2020 a Mayo 2023 dando 269 observaciones. 2. De “Precios-Promedio-Nacionales-Diarios-2023” se importaron los precios promedio nacionales de forma diaria de las columnas de Super, Regular, Diesel y Gas Licuado Propano de 25 lbs un total de 942 observaciones. 3. De “CONSUMO-2023-05” se importaron los volúmenes de consumo mensuales de Diesel bajo zufre, Gasolina regulars, Gasolinea superior y Diesel alto azufre un total de 281 observaciones.

0.2 Analisis exploratorio

```
[27]: plt.figure(figsize=(13,6))
      df_imp['Gas licuado de petróleo'].plot(color='purple', linewidth=0.8)
      df_imp['Gasolina regular'].plot(color='red', linewidth=0.8)
      df_imp['Gasolina superior'].plot(color='orange', linewidth=0.8)
      df_imp['Diesel'].plot(color='blue', linewidth=0.5)
      plt.show()
```

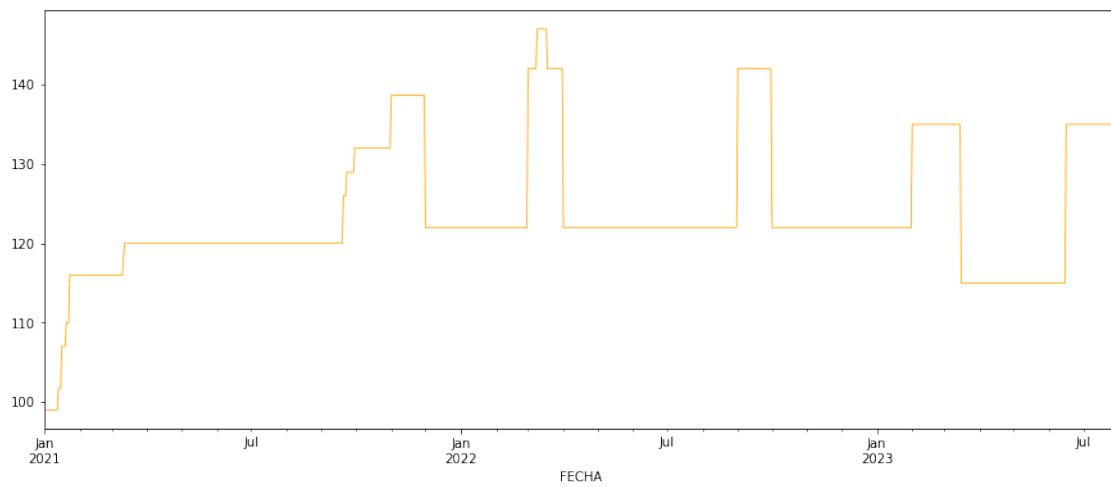



```
[28]: sns.pairplot(df_imp)
plt.show()
```

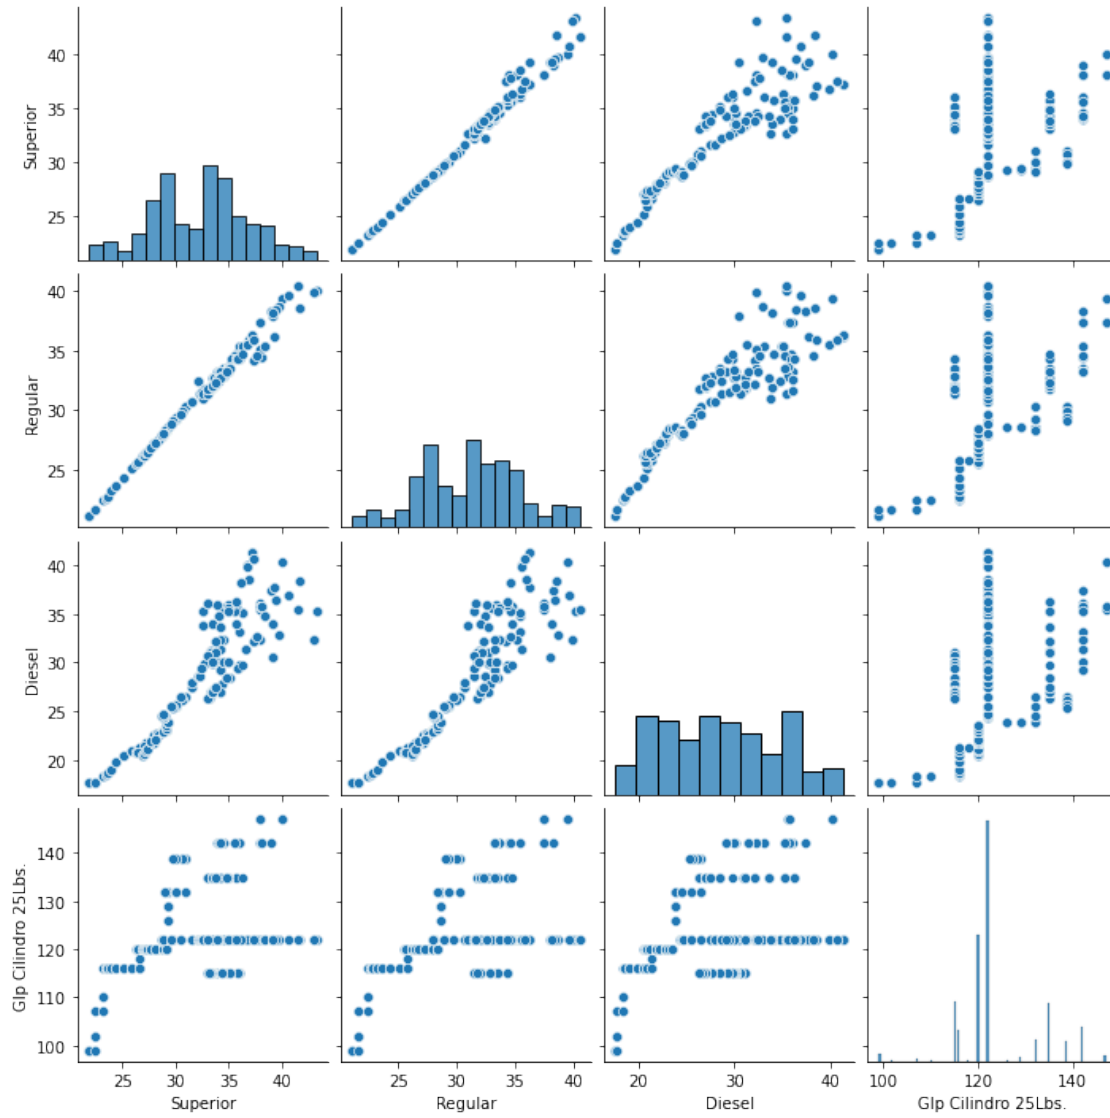



```
[29]: plt.figure(figsize=(15,6))
df_precios['Superior'].plot(color='purple', linewidth=0.8)
df_precios['Regular'].plot(color='red', linewidth=0.8)
df_precios['Diesel'].plot(color='blue', linewidth=0.5)
plt.show()


plt.figure(figsize=(15,6))
df_precios['Glp Cilindro 25Lbs.'].plot(color='orange', linewidth=0.8)
plt.show()
```



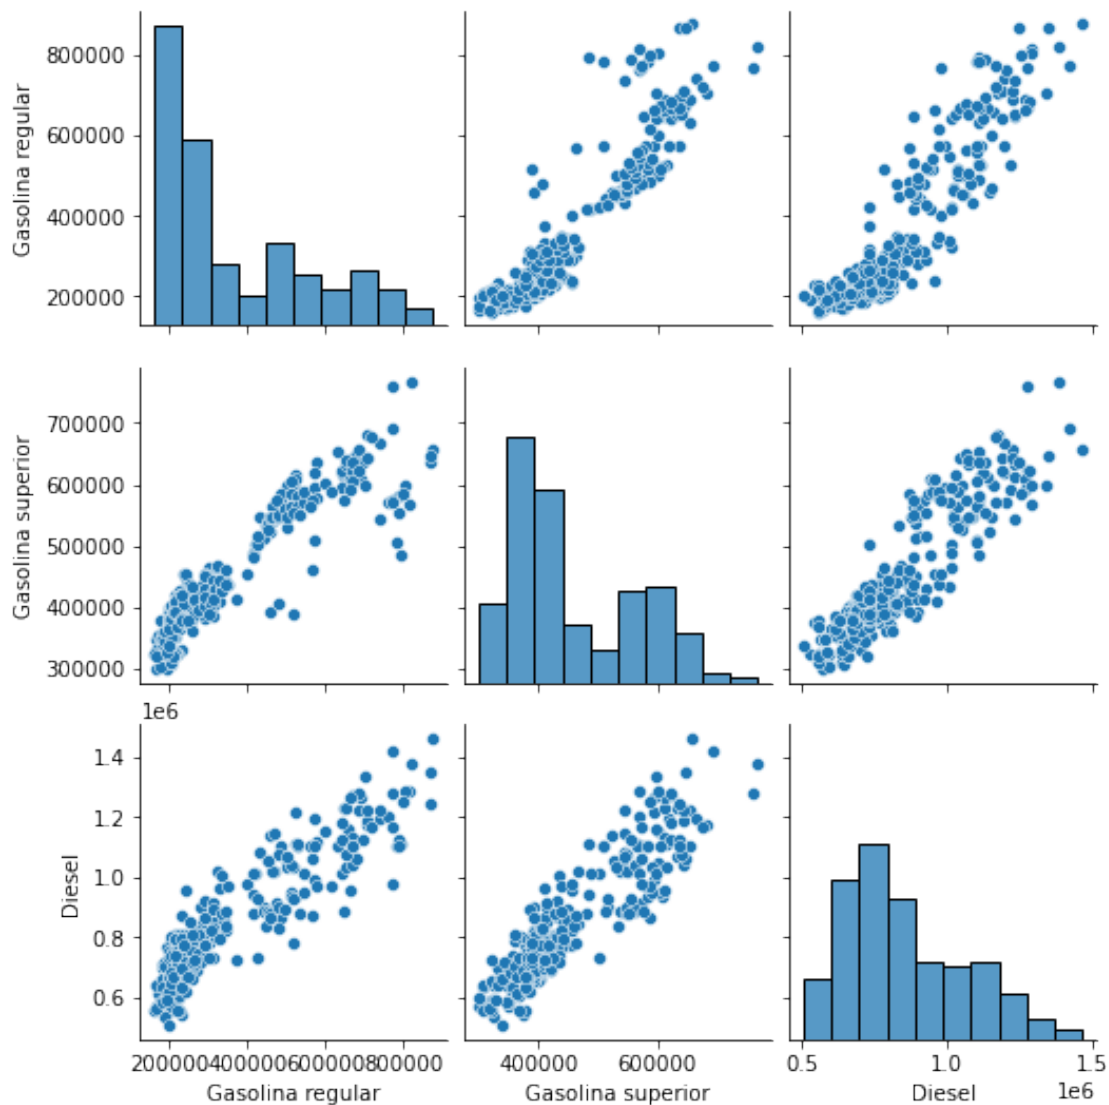
```
[30]: sns.pairplot(df_precios)
plt.show()
```



```
[31]: plt.figure(figsize=(15,6))
df_consumo['Gasolina regular'].plot(color='red', linewidth=0.8)
df_consumo['Gasolina superior'].plot(color='orange', linewidth=0.8)
df_consumo['Diesel'].plot(color='blue', linewidth=0.5)
plt.show()
```



```
[32]: sns.pairplot(df_consumo)
plt.show()
```



1 ARIMA

Hemo decido utilizar las siguientes columnas para trabajar con arima: - Importacion: Gas licuado de petróleo como df1 - Precios: Gasolina superior como df2 - Consumo: Diesel df3

```
[33]: df1 = df_imp['Gasolina regular']
      df2 = df_precios['Superior']
      df3 = df_consumo['Diesel']
```


```
[34]: df1.head(), df2.head(), df3.head()
```

```
[34]: (Fecha
      2001-01-01    177776.50
      2001-02-01    123115.99
      2001-03-01    161726.42
      2001-04-01    127338.74
      2001-05-01    168730.19
      Name: Gasolina regular, dtype: float64,
      FECHA
      2021-01-01     21.91
      2021-01-02     21.91
      2021-01-03     21.91
      2021-01-04     21.91
      2021-01-05     21.91
      Name: Superior, dtype: float64,
      Fecha
      2000-01-01    634667.06
      2000-02-01    642380.66
      2000-03-01    699807.25
      2000-04-01    586803.98
      2000-05-01    656948.20
      Name: Diesel, dtype: float64)
```


1.0.1 Importacion de Gasolina regular

Exploracion

```
[35]: plt.figure(figsize=(15,6))
      plt.plot(df1)
      plt.gca().set(title="Volumen de Importación", xlabel="Mes-Año",
      ↪ylabel="Cantidad de Galones")
      plt.show()
```



```
[36]: # Se calcula la media móvil y la desviación estandar móvil de los últimos 12
      ↪ meses.
mediaMovil = df1.rolling(window=12).mean()
deMovil = df1.rolling(window=12).std()
# Se grafican los resultados.
plt.figure(figsize=(15,6))
original = plt.plot(df1, color="blue", label="Original")
media = plt.plot(mediaMovil, color='red', label = 'Media Movil')
ds = plt.plot(deMovil,color='black', label = 'Desviación Estándar Móvil')
plt.legend(loc = 'best')
plt.title('Media y desviación estándar móvil')
plt.show(block=False)
```



Se observa que ni la media ni la varianza son estacionarias

```
[37]: result = seasonal_decompose(df1)
      # Create a new figure with a specific size
fig, (ax1, ax2, ax3, ax4) = plt.subplots(4, 1, figsize=(12, 6))


# Plot the components on separate subplots
result.observed.plot(ax=ax1)
ax1.set_ylabel('Observed')

result.trend.plot(ax=ax2)
ax2.set_ylabel('Trend')

result.seasonal.plot(ax=ax3)
ax3.set_ylabel('Seasonal')


result.resid.plot(ax=ax4)
ax4.set_ylabel('Residual')
```

```
# Adjust layout and display the plot
plt.tight_layout()
plt.show()
```



```
[38]: df1_log = np.log(df1)
plt.plot(df1_log)
```

```
[38]: [<matplotlib.lines.Line2D at 0x228e1540be0>]
```




```
[39]: mediaMovil_log = df1_log.rolling(window=12).mean()
deMovil_log = df1_log.rolling(window=12).std()
#Graficando los resultados con logaritmo.
#serie_df1_log = plt.plot(df1_log, color="blue", label="df1_log")
#media_movil_df1_log = plt.plot(mediaMovil_log, color='red', label = 'Media_
↳ Movil_log')
ds_log = plt.plot(deMovil_log,color='black', label = 'Desviación Estándar_
↳ Móvil_log')
plt.legend(loc = 'best')
plt.title('Desviación estándar móvil transformada con LN')
plt.show(block=False)
```



```
[40]: print('\033[1mResultados del Test de Dickey Fuller de la serie Importación_
↳ Gasolina Regular\033[0m')
df1_test = adfuller(df1, autolag='AIC')
salida_df1 = pd.Series(df1_test[0:4], index=['Estadístico de_
↳ prueba', 'p-value', '# de retardos usados', '# de observaciones usadas'])
for key,value in df1_test[4].items():
    salida_df1['Critical Value (%s)'%key] = value
print(salida_df1)
```

Resultados del Test de Dickey Fuller de la serie Importación Gasolina

Regular

Estadístico de prueba	1.090748
p-value	0.995137
# de retardos usados	9.000000
# de observaciones usadas	259.000000
Critical Value (1%)	-3.455853
Critical Value (5%)	-2.872765
Critical Value (10%)	-2.572752

dtype: float64

```
[41]: print('\033[1mResultados del Test de Dickey Fuller para 1 diferenciación de la_
      ↪serie Importación Gasolina regular\033[0m')
df1_diff = df1.diff()
df1_diff.dropna(inplace=True)
df1_test = adfuller(df1_diff)
salida_df1 = pd.Series(df1_test[0:4], index=['Estadístico de_
      ↪prueba', 'p-value', '# de retardos usados', '# de observaciones usadas'])
for key,value in df1_test[4].items():
    salida_df1['Critical Value (%s)'%key] = value
print(salida_df1)
```

Resultados del Test de Dickey Fuller para 1 diferenciación de la serie


Importación Gasolina regular

Estadístico de prueba	-9.852868e+00
p-value	4.437791e-17
# de retardos usados	8.000000e+00
# de observaciones usadas	2.590000e+02
Critical Value (1%)	-3.455853e+00
Critical Value (5%)	-2.872765e+00
Critical Value (10%)	-2.572752e+00

dtype: float64


```
[42]: plt.figure(figsize=(12, 6))
      plt.title('Importaciones de Gasolina Regular con diferenciación')
      plt.plot(df1_diff)
```

```
[42]: [<matplotlib.lines.Line2D at 0x228ddaa31f0>]
```



```
[43]: plt.figure(figsize=(12, 4))
df1_log_diff = df1_log.diff()
plt.title('Importaciones de Gasolina Regular con Transformación logarítmica y
↪diferenciación')
plt.plot(df1_log_diff)
```

[43]: [<matplotlib.lines.Line2D at 0x228e05ff580>]



```
[44]: df1_log_diff.dropna(inplace = True)
df1_log_diff_acf = acf(df1_log_diff,nlags=5,fft=False)
df1_log_diff_pacf = pacf(df1_log_diff, nlags=36)
df1_log_diff_acf, df1_log_diff_pacf
```


```
[44]: (array([ 1.          , -0.63980275,  0.22972522, -0.1397927 ,  0.05508003,
              0.03811285]),
       array([ 1.          , -0.64219901, -0.30798694, -0.27215995, -0.25438727,
              -0.10156584, -0.04916746, -0.2515719 , -0.109741  , -0.19573169,
              -0.12725689, -0.06880904,  0.02507143, -0.17524294, -0.09741736,
              -0.05721368,  0.00913106,  0.10158407, -0.15873323, -0.15162316,
              -0.03222279, -0.0305563 ,  0.00159169, -0.04679743, -0.09825368,
              -0.12212896, -0.01962758,  0.01365864,  0.05888865, -0.03679905,
              0.09447004, -0.06584161,  0.02074825,  0.03536952, -0.24681957,
              0.1588869 ,  0.03900979]))
```

```
[45]: plt.rcParams['figure.figsize'] = [14, 5]
plt.rcParams['figure.dpi'] = 100 # 200 e.g. is really fine, but slower

#Plot ACF:
plt.subplot(121)
plt.plot(acf(df1_log_diff,nlags=12,fft=False))
plt.axhline(y=0,linestyle='--',color='gray')
plt.axhline(y=-1.96/np.sqrt(len(df1_log_diff)),linestyle='--',color='gray')
plt.axhline(y=1.96/np.sqrt(len(df1_log_diff)),linestyle='--',color='gray')
plt.title('Función de Autocorrelación con 12 retardos')

plt.subplot(122)
plt.plot(acf(df1_log_diff,nlags=24,fft=False))
plt.axhline(y=0,linestyle='--',color='gray')
plt.axhline(y=-1.96/np.sqrt(len(df1_log_diff)),linestyle='--',color='gray')
plt.axhline(y=1.96/np.sqrt(len(df1_log_diff)),linestyle='--',color='gray')
plt.title('Función de Autocorrelación con 24 retardos')

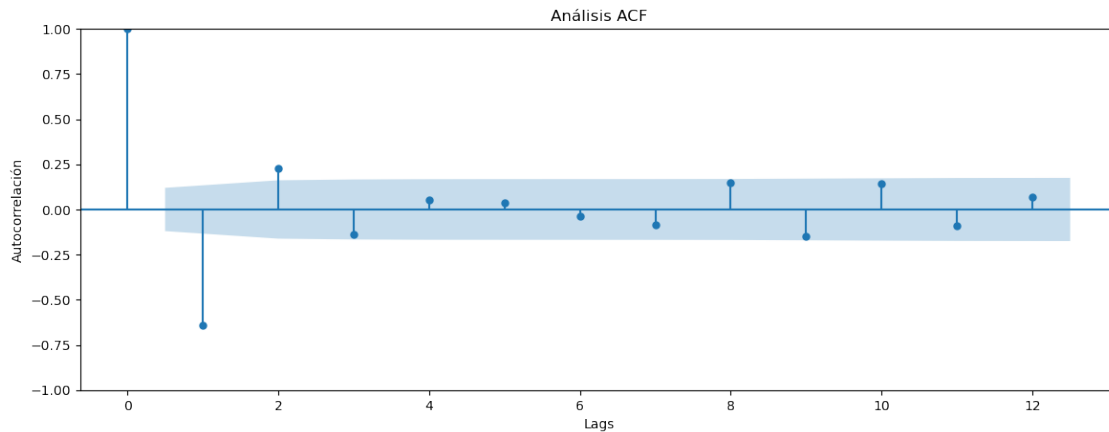
plt.tight_layout()
```



```
[46]: plt.figure(figsize=(10, 6))
plot_acf(df1_log_diff, lags=12)
```

```
plt.xlabel('Lags')
plt.ylabel('Autocorrelación')
plt.title('Análisis ACF')
plt.show()
```

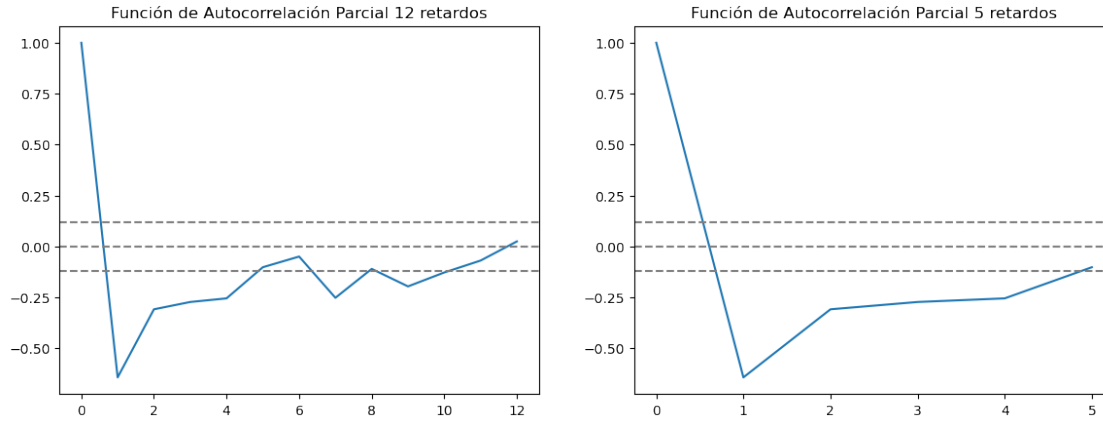
<Figure size 1000x600 with 0 Axes>



```
[47]: #plot PACF
plt.subplot(121)
plt.plot(pacf(df1_log_diff, nlags=12))
plt.axhline(y=0,linestyle='--',color='gray')
plt.axhline(y=-1.96/np.sqrt(len(df1_log_diff)),linestyle='--',color='gray')
plt.axhline(y=1.96/np.sqrt(len(df1_log_diff)),linestyle='--',color='gray')
plt.title('Función de Autocorrelación Parcial 12 retardos')

plt.subplot(122)
plt.plot(pacf(df1_log_diff, nlags=5))
plt.axhline(y=0,linestyle='--',color='gray')
plt.axhline(y=-1.96/np.sqrt(len(df1_log_diff)),linestyle='--',color='gray')
plt.axhline(y=1.96/np.sqrt(len(df1_log_diff)),linestyle='--',color='gray')
plt.title('Función de Autocorrelación Parcial 5 retardos')
```

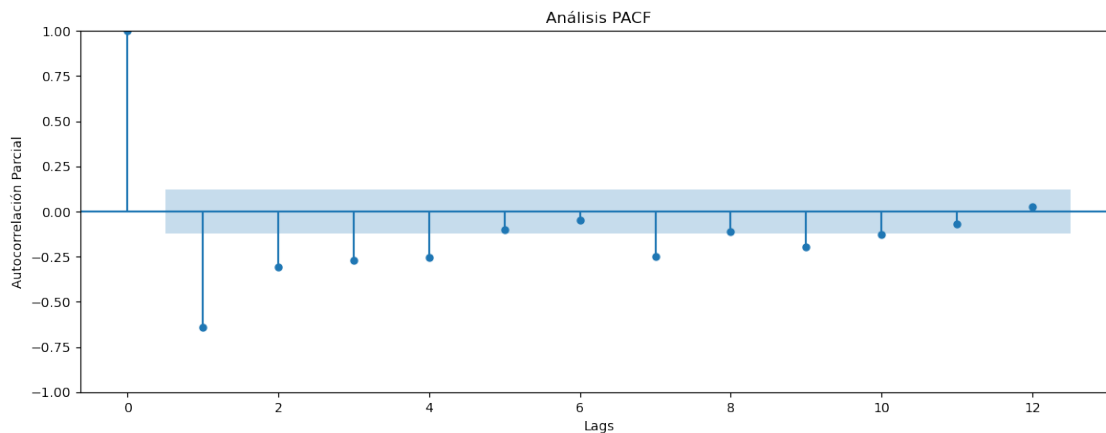
```
[47]: Text(0.5, 1.0, 'Función de Autocorrelación Parcial 5 retardos')
```



```
[48]: plt.figure(figsize=(10, 6))
plot_pacf(df1_log_diff, lags=12)
plt.xlabel('Lags')
plt.ylabel('Autocorrelación Parcial')
plt.title('Análisis PACF')
plt.show()
```

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\graphics\tsaplots.py:348:
FutureWarning: The default method 'yw' can produce PACF values outside of the
[-1,1] interval. After 0.13, the default will change to unadjusted Yule-Walker
('ywm'). You can use this method now by setting method='ywm'.
warnings.warn(

<Figure size 1000x600 with 0 Axes>



para el P podemos usar 5

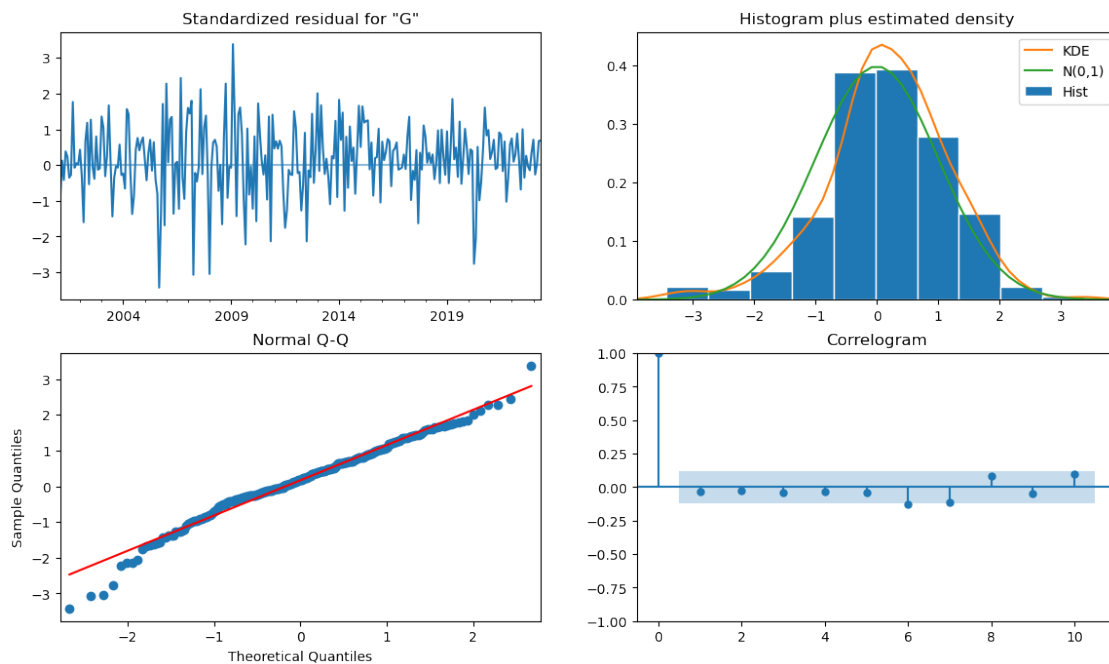
No parece haber estacionalidad

```
[49]: #df1_log = df1_log.resample("M").last()
df1_modelo512 = ARIMA(df1_log, order=(5,1,2), freq='MS')
resultado_df1_modelo512 = df1_modelo512.fit()
print(resultado_df1_modelo512.summary().tables[1])
```

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.
self._init_dates(dates, freq)

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	-0.4176	1.127	-0.370	0.711	-2.627	1.792
ar.L2	0.0011	0.336	0.003	0.997	-0.657	0.660
ar.L3	-0.0622	0.111	-0.561	0.575	-0.279	0.155
ar.L4	-0.0107	0.155	-0.069	0.945	-0.314	0.292
ar.L5	0.0508	0.090	0.564	0.573	-0.126	0.228
ma.L1	-0.6447	1.128	-0.572	0.568	-2.855	1.566
ma.L2	-0.1559	0.889	-0.175	0.861	-1.898	1.586
sigma2	0.0711	0.006	12.781	0.000	0.060	0.082

```
[50]: resultado_df1_modelo512.plot_diagnostics(figsize=(14, 8))
plt.show()
```



Los residuos presentan una distribución normal, y parecen comportarse como un ruido blanco dado que en el correlograma no hay autocorrelaciones significativas

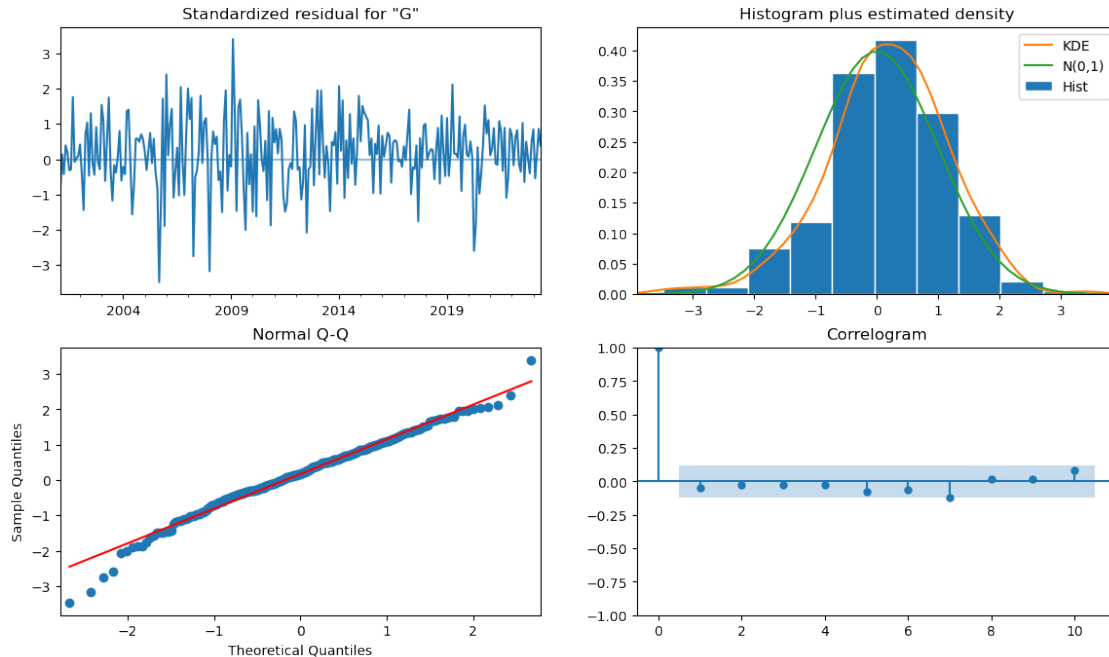
```
[51]: df1_modelo513 = ARIMA(df1_log, order=(5,1,3), freq='MS')
      resultado_df1_modelo513 = df1_modelo513.fit()
      print(resultado_df1_modelo513.summary().tables[1])
```

```
C:\ProgramData\Anaconda3\lib\site-
packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency
information was provided, so inferred frequency MS will be used.
    self._init_dates(dates, freq)
C:\ProgramData\Anaconda3\lib\site-
packages\statsmodels\tsa\statespace\sarimax.py:966: UserWarning: Non-stationary
starting autoregressive parameters found. Using zeros as starting parameters.
    warn('Non-stationary starting autoregressive parameters')
C:\ProgramData\Anaconda3\lib\site-
packages\statsmodels\tsa\statespace\sarimax.py:978: UserWarning: Non-invertible
starting MA parameters found. Using zeros as starting parameters.
    warn('Non-invertible starting MA parameters found.')
```

```
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1          -1.1787      0.098     -12.034      0.000      -1.371      -0.987
ar.L2          -1.0688      0.159      -6.703      0.000      -1.381      -0.756
ar.L3          -0.2446      0.189      -1.295      0.195      -0.615       0.125
ar.L4          -0.0549      0.137      -0.401      0.688      -0.323       0.213
ar.L5          -0.0112      0.078      -0.143      0.886      -0.164       0.142
ma.L1           0.1467      0.084       1.752      0.080      -0.017       0.311
ma.L2           0.1211      0.078       1.549      0.121      -0.032       0.274
ma.L3          -0.7740      0.069     -11.251      0.000      -0.909      -0.639
sigma2          0.0693      0.005      12.638      0.000       0.059       0.080
=====
```

```
C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\base\model.py:604:
ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check
mle_retvals
    warnings.warn("Maximum Likelihood optimization failed to "
```

```
[52]: resultado_df1_modelo513.plot_diagnostics(figsize=(14, 8))
      plt.show()
```

```
[53]: print("Resultados de AIC (Akaike information criterion)")
print("Modelo 512=",resultado_df1_modelo512.aic)
print("Modelo 513=",resultado_df1_modelo513.aic)
print("Resultados de BIC (Bayesian information criterion)")
print("Modelo 512=",resultado_df1_modelo512.bic)
print("Modelo 513=",resultado_df1_modelo513.bic)
```

```
Resultados de AIC (Akaike information criterion)
Modelo 512= 69.58469079482742
Modelo 513= 63.007380612464054
Resultados de BIC (Bayesian information criterion)
Modelo 512= 98.31258663891427
Modelo 513= 95.32626343706175
```

De acuerdo a ambos indicadores es mejor el modelo $p=5$, $d=1$, $q=3$ por lo que este es el que será usado para predecir pues tiene un valor menor en AIC y en BIC Predicción Importaciones de Gasolina Regular con el mejor modelo

```
[54]: # Generate forecast for the next 12 months
forecast_steps = 12
forecast = resultado_df1_modelo513.forecast(steps=forecast_steps)
forecast_original_scale = np.exp(forecast)

# Create a date range for the forecasted months
last_date = df1.index[-6] # Assuming 'df' is your original data DataFrame
```


```

forecast_index = pd.date_range(start=last_date, periods=forecast_steps + 1,
    ↪freq='M')[1:]

# Create a DataFrame to hold the forecasted values
forecast_df = pd.DataFrame({'Forecast': forecast}, index=forecast_index)

# Plot the original data and the forecasted values
plt.figure(figsize=(10, 6))
plt.plot(df1, label='Original Data')
plt.plot(forecast_df.index, forecast_original_scale, label='Forecast',
    ↪color='orange')
plt.xlabel('Date')
plt.ylabel('Value')
plt.title('ARIMA para 36 meses')
plt.legend()
plt.show()

```



1.0.2 Precios de Gasolina Superior


Exploracion

```


[55]: plt.figure(figsize=(15,6))
      plt.plot(df2)

```

```
plt.gca().set(title="Precios de Gasolina Superior", xlabel="Mes-Año",
    ↪ylabel="Precio")
plt.show()
```



```
[56]: # Se calcula la media móvil y la desviación estandar móvil de los últimos 12
    ↪meses.
mediaMovil_df2 = df2.rolling(window=7).mean()
deMovil_df2 = df2.rolling(window=7).std()
# Se grafican los resultados.
plt.figure(figsize=(15,6))
original = plt.plot(df2, color="blue", label="Original")
media = plt.plot(mediaMovil_df2, color='red', label = 'Media Movil')
ds = plt.plot(deMovil_df2,color='black', label = 'Desviación Estándar Móvil')
plt.legend(loc = 'best')
plt.title('Media y desviación estándar móvil')
plt.show(block=False)
```



```
[57]: result = seasonal_decompose(df2)
# Create a new figure with a specific size
fig, (ax1, ax2, ax3, ax4) = plt.subplots(4, 1, figsize=(12, 6))


# Plot the components on separate subplots
result.observed.plot(ax=ax1)
ax1.set_ylabel('Observed')

result.trend.plot(ax=ax2)
ax2.set_ylabel('Trend')

result.seasonal.plot(ax=ax3)
ax3.set_ylabel('Seasonal')

result.resid.plot(ax=ax4)
ax4.set_ylabel('Residual')

# Adjust layout and display the plot
plt.tight_layout()
plt.show()
```



No parece haber tendencia

```
[58]: print('\033[1mResultados del Test de Dickey Fuller de la serie Precios Gasolina_↵Superior\033[0m')
df2_test = adfuller(df2, autolag='AIC')
```

```

salida_df2 = pd.Series(df2_test[0:4], index=['Estadístico de
↳prueba', 'p-value', '# de retardos usados', '# de observaciones usadas'])
for key,value in df2_test[4].items():
    salida_df2['Critical Value (%s)'%key] = value
print(salida_df2)

```

Resultados del Test de Dickey Fuller de la serie Precios Gasolina

Superior

Estadístico de prueba	-2.345554
p-value	0.157696
# de retardos usados	7.000000
# de observaciones usadas	934.000000
Critical Value (1%)	-3.437371
Critical Value (5%)	-2.864639
Critical Value (10%)	-2.568420
dtype:	float64

El valor P aún es mayor de 0.05, no se puede rechazar

```

[59]: print('\033[1mResultados del Test de Dickey Fuller para 1 diferenciación de la
↳serie Precios Gasolina Superior\033[0m')
df2_diff = df2.diff()
df2_diff.dropna(inplace=True)
df2_test = adfuller(df2_diff)
salida_df2 = pd.Series(df2_test[0:4], index=['Estadístico de
↳prueba', 'p-value', '# de retardos usados', '# de observaciones usadas'])
for key,value in df2_test[4].items():
    salida_df2['Critical Value (%s)'%key] = value
print(salida_df2)

```

Resultados del Test de Dickey Fuller para 1 diferenciación de la serie

Precios Gasolina Superior

Estadístico de prueba	-8.535789e+00
p-value	1.008898e-13
# de retardos usados	6.000000e+00
# de observaciones usadas	9.340000e+02
Critical Value (1%)	-3.437371e+00
Critical Value (5%)	-2.864639e+00
Critical Value (10%)	-2.568420e+00
dtype:	float64

Con una diferenciación se anula acepta la Hipótesis alternativa

```


[60]: plt.plot(df2_diff)

```

```

[60]: [<matplotlib.lines.Line2D at 0x228e1888040>]

```




```
[61]: df2_diff.dropna(inplace = True)
df2_diff_acf = acf(df2_diff,nlags=12,fft=False)
df2_diff_pacf = pacf(df2_diff, nlags=36)
df2_diff_acf, df2_diff_pacf
```

```
[61]: (array([ 1.00000000e+00, -1.68259614e-03,  3.24949192e-02,  4.64532552e-03,
 1.48948761e-02, -1.49575012e-03,  2.99319549e-02,  2.34814760e-01,
 5.49055449e-03,  5.82181753e-03, -1.41412108e-02,  4.27028945e-04,
-2.37006595e-03]),
array([ 1.00000000e+00, -1.68438613e-03,  3.25613862e-02,  4.77316740e-03,
 1.39293581e-02, -1.76074392e-03,  2.91991345e-02,  2.37155701e-01,
 6.29794789e-03, -9.94653598e-03, -1.80803917e-02, -6.17082383e-03,
 1.33502459e-04, -7.28867855e-03, -9.56818921e-02, -1.41073124e-02,
 2.71259281e-02, -6.15290964e-03,  2.25098490e-03,  3.51999768e-03,
-4.00348411e-02, -2.02119141e-02,  2.26717790e-03, -5.72068553e-04,
 3.16591770e-02, -1.96180529e-03,  1.46151740e-03, -2.85567143e-02,
-3.92765859e-02,  3.33340379e-03,  9.99064945e-03, -2.78959076e-02,
-3.60200750e-03,  3.60790103e-03, -4.26500646e-02, -5.52877200e-02,
-4.98884312e-03]))
```

```
[62]: plt.figure(figsize=(10, 6))
plot_acf(df2_diff, lags=36)
plt.xlabel('Lags')
plt.ylabel('Autocorrelación')
plt.title('Análisis ACF "Precios Gasolina Superior"')
plt.show()
```

<Figure size 1000x600 with 0 Axes>



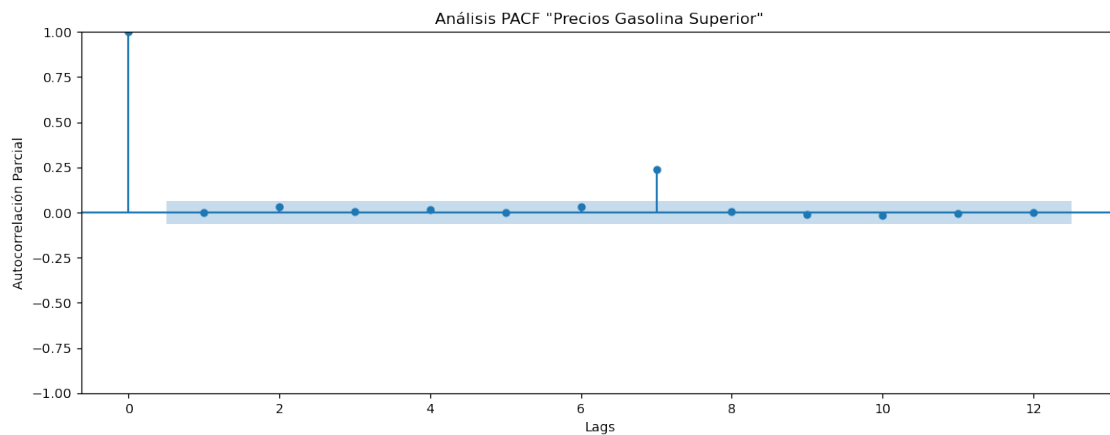
Se sugiere un valor de desde de 7 para el Q No Parece haber un componente de temporadas

```
[63]: plt.figure(figsize=(8, 6))
plot_pacf(df2_diff, lags=12)
plt.xlabel('Lags')
plt.ylabel('Autocorrelación Parcial')
plt.title('Análisis PACF "Precios Gasolina Superior"')
plt.show()
```

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\graphics\tsaplots.py:348:
FutureWarning: The default method 'yw' can produce PACF values outside of the
[-1,1] interval. After 0.13, the default will change to unadjusted Yule-Walker
(`'yw'`). You can use this method now by setting `method='ywm'`.

warnings.warn(

<Figure size 800x600 with 0 Axes>



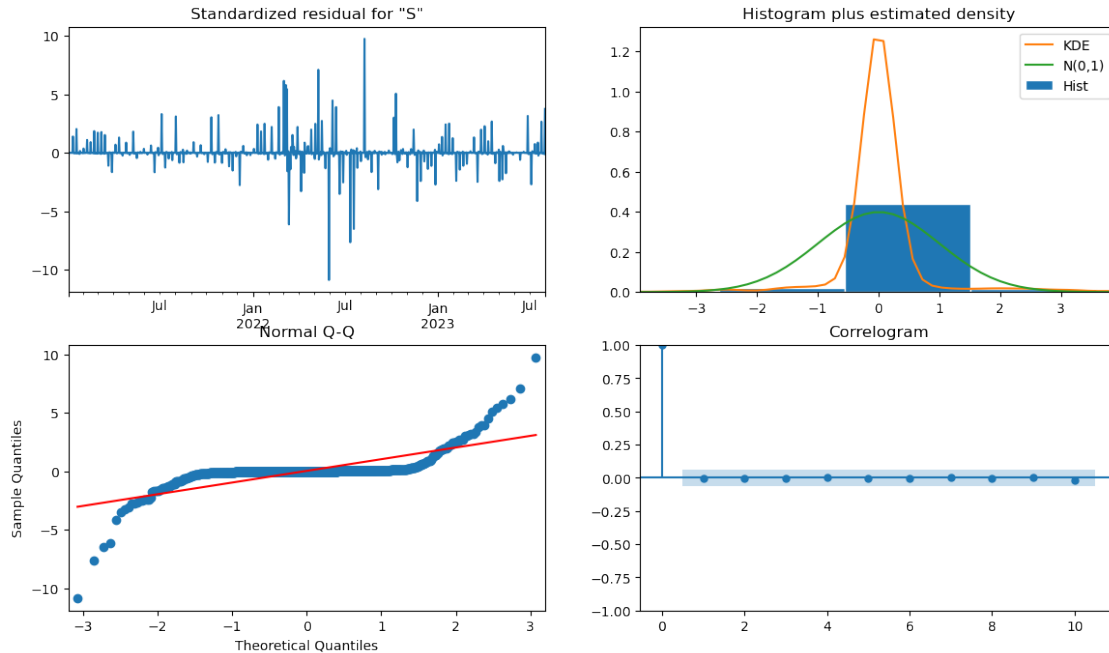
Se sugiere un componente de 7 para el P

```
[64]: df2_modelo717 = ARIMA(df2, order=(7,1,7), freq='D')
      resultado_df2_modelo717 = df2_modelo717.fit()
      print(resultado_df2_modelo717.summary().tables[1])
```

```
C:\ProgramData\Anaconda3\lib\site-
packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency
information was provided, so inferred frequency D will be used.
    self._init_dates(dates, freq)
```

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.0252	0.470	0.054	0.957	-0.896	0.946
ar.L2	-0.0077	0.461	-0.017	0.987	-0.912	0.896
ar.L3	0.0285	0.189	0.151	0.880	-0.342	0.399
ar.L4	-0.0277	0.370	-0.075	0.940	-0.754	0.698
ar.L5	-0.0141	0.616	-0.023	0.982	-1.221	1.193
ar.L6	0.1327	0.231	0.575	0.565	-0.319	0.585
ar.L7	-0.0736	0.098	-0.752	0.452	-0.265	0.118
ma.L1	-0.0350	0.422	-0.083	0.934	-0.863	0.793
ma.L2	0.0456	0.433	0.105	0.916	-0.803	0.895
ma.L3	-0.0240	0.157	-0.153	0.878	-0.331	0.283
ma.L4	0.0422	0.376	0.112	0.911	-0.695	0.779
ma.L5	0.0042	0.556	0.007	0.994	-1.085	1.093
ma.L6	-0.0996	0.218	-0.458	0.647	-0.526	0.327
ma.L7	0.3337	0.102	3.274	0.001	0.134	0.533
sigma2	0.1162	0.001	89.942	0.000	0.114	0.119

```
[65]: resultado_df2_modelo717.plot_diagnostics(figsize=(14, 8))
      plt.show()
```

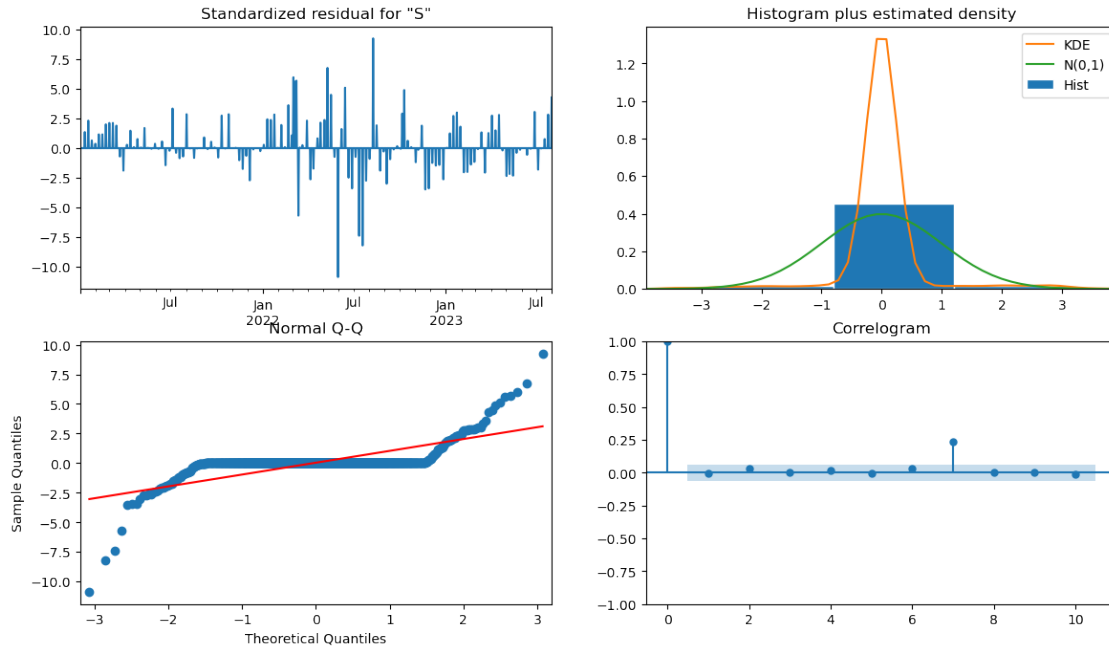



```
[66]: df2_modelo010 = ARIMA(df2, order=(0,1,0), freq='D')
      resultado_df2_modelo010 = df2_modelo010.fit()
      print(resultado_df2_modelo010.summary().tables[1])
```

```
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
sigma2          0.1246      0.001     98.495      0.000      0.122      0.127
=====
```

```
C:\ProgramData\Anaconda3\lib\site-
packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency
information was provided, so inferred frequency D will be used.
    self._init_dates(dates, freq)
```

```
[67]: resultado_df2_modelo010.plot_diagnostics(figsize=(14, 8))
      plt.show()
```



```
[68]: df2_modelo111 = ARIMA(df2, order=(1,1,1), freq='D')
      resultado_df2_modelo111 = df2_modelo111.fit()
      print(resultado_df2_modelo111.summary().tables[1])
```

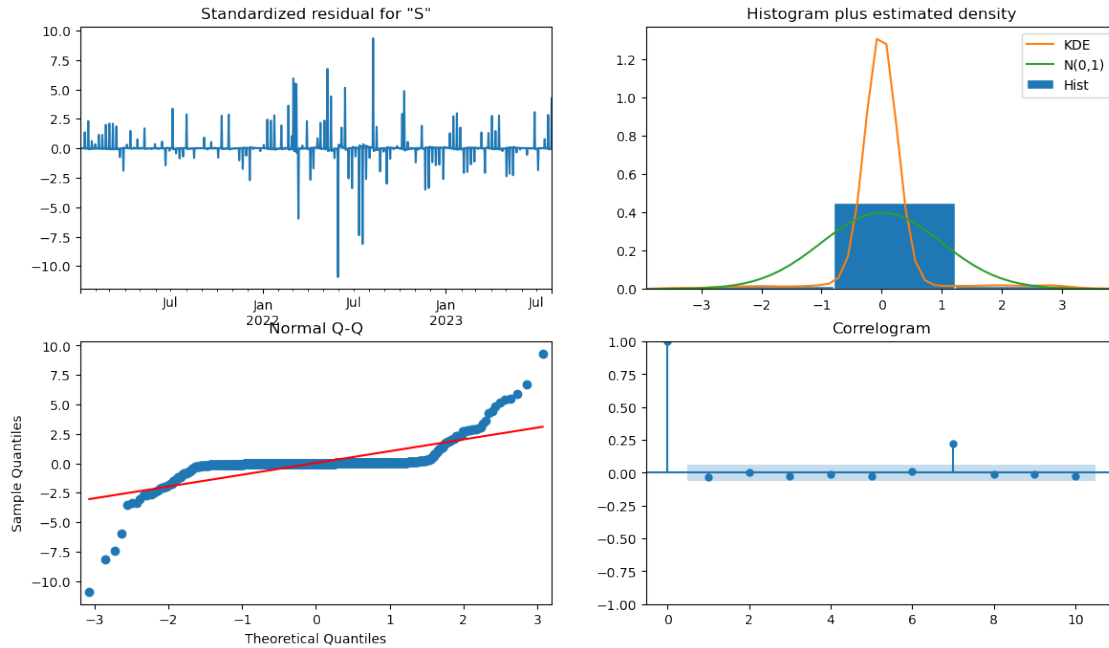
```
C:\ProgramData\Anaconda3\lib\site-
packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency
information was provided, so inferred frequency D will be used.
      self._init_dates(dates, freq)
```

```
=====
```

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.9017	0.107	8.460	0.000	0.693	1.111
ma.L1	-0.8717	0.133	-6.552	0.000	-1.132	-0.611
sigma2	0.1240	0.002	80.865	0.000	0.121	0.127

```
=====
```

```
[69]: resultado_df2_modelo111.plot_diagnostics(figsize=(14, 8))
      plt.show()
```



```
[70]: print("Resultados de AIC (Akaike information criterion)")
print("Modelo 717=", resultado_df2_modelo717.aic)
print("Modelo 010=", resultado_df2_modelo010.aic)
print("Modelo 111=", resultado_df2_modelo111.aic)
print("Resultados de BIC (Bayesian information criterion)")
print("Modelo 717=", resultado_df2_modelo717.bic)
print("Modelo 010=", resultado_df2_modelo010.bic)
print("Modelo 111=", resultado_df2_modelo111.bic)
```

```
Resultados de AIC (Akaike information criterion)
Modelo 717= 675.8725376274368
Modelo 010= 712.6509335511615
Modelo 111= 712.2318202601184
Resultados de BIC (Bayesian information criterion)
Modelo 717= 748.5766847212175
Modelo 010= 717.4978766907469
Modelo 111= 726.7726496788746
```

el mejor modelo para ser el 010 procedemos a pronosticar ##### Predicción con el mejor modelo
Precios Gasolina Superior

```
[71]: # Generate forecast for the next 90 dias
forecast_steps = 365
forecast = resultado_df2_modelo010.forecast(steps=forecast_steps)


# Create a date range for the forecasted months
last_date = df2[-1] # Assuming 'df' is your original data DataFrame
```

```

forecast_index = pd.date_range(start=last_date, periods=forecast_steps + 1,
    ↪freq='D')[1:]

# Plot the original data and the forecasted values
plt.figure(figsize=(10, 6))
plt.plot(df2, label='Original Data')
plt.plot(forecast, label='Forecast', color='orange')
plt.xlabel('Date')
plt.ylabel('Value')
plt.title('ARIMA para 90 dias Precios Gasolina Superior')
plt.legend()
plt.show()

```




1.0.3 Consumos Gasolina Diesel

Exploración

```


[72]: plt.figure(figsize=(15,6))
plt.plot(df3)
plt.gca().set(title="Volumen de Consumo", xlabel="Mes-Año", ylabel="Diesel")
plt.show()

```



[73]: *# Se calcula la media móvil y la desviación estandar móvil de los últimos 12 meses.*

```
mediaMovil_df3 = df3.rolling(window=12).mean()
deMovil_df3 = df3.rolling(window=12).std()
# Se grafican los resultados.
plt.figure(figsize=(12,6))
original_df3 = plt.plot(df3, color="blue", label="Original")
media_df3 = plt.plot(mediaMovil_df3, color='red', label = 'Media Movil')
ds_df3 = plt.plot(deMovil_df3,color='black', label = 'Desviación Estándar Móvil')
plt.legend(loc = 'best')
plt.title('Media y desviación estándar móvil')
plt.show(block=False)
```



```
[74]: result_df3 = seasonal_decompose(df3)
# Create a new figure with a specific size
fig_df3, (ax1, ax2, ax3, ax4) = plt.subplots(4, 1, figsize=(12, 6))


# Plot the components on separate subplots
result.observed.plot(ax=ax1)
ax1.set_ylabel('Observed')

result.trend.plot(ax=ax2)
ax2.set_ylabel('Trend')

result.seasonal.plot(ax=ax3)
ax3.set_ylabel('Seasonal')


result.resid.plot(ax=ax4)
ax4.set_ylabel('Residual')

# Adjust layout and display the plot
plt.tight_layout()
plt.show()
```




```
[75]: df3_log = np.log(df3)
plt.plot(df3_log)
```

```
[75]: [<matplotlib.lines.Line2D at 0x228e02e5700>]
```



```
[76]: mediaMovil_log_df3 = df3_log.rolling(window=12).mean()
deMovil_log_df3 = df3_log.rolling(window=12).std()
#Graficando los resultados con logaritmo.
#serie_df3_log = plt.plot(df3_log, color="blue", label="df3_log")
#media_movil_df3_log = plt.plot(mediaMovil_log_df3, color='red', label = 'Media_
↳Movil_log')
ds_log = plt.plot(deMovil_log_df3,color='black', label = 'Desviación Estándar_
↳Móvil_log')
plt.legend(loc = 'best')
plt.title('Desviación estándar móvil Diesel transformada con LN')
plt.show(block=False)
```



Se tienen dos claros picos de la crisis del 2008 y el Covid pero en general es bastante estable

```
[77]: print('\033[1mResultados del Test de Dickey Fuller de la serie Consumo_
↳Diesel\033[0m')
df3_test = adfuller(df3, autolag='AIC')
```

```

salida_df3 = pd.Series(df3_test[0:4], index=['Estadístico de prueba', 'p-value', '# de retardos usados', '# de observaciones usadas'])
for key,value in df3_test[4].items():
    salida_df3['Critical Value (%s)'%key] = value
print(salida_df3)

```

Resultados del Test de Dickey Fuller de la serie Consumo Diesel

Estadístico de prueba	0.018032
p-value	0.959998
# de retardos usados	12.000000
# de observaciones usadas	268.000000
Critical Value (1%)	-3.454988
Critical Value (5%)	-2.872386
Critical Value (10%)	-2.572549

dtype: float64

El valor p es mayor a 0.05 se procede a diferenciar 1 vez

```

[78]: print('\033[1mResultados del Test de Dickey Fuller para 1 diferenciación de la serie Consumo Diesel\033[0m')
df3_diff = df3.diff()
df3_diff.dropna(inplace=True)
df3_test = adfuller(df3_diff)
salida_df3 = pd.Series(df3_test[0:4], index=['Estadístico de prueba', 'p-value', '# de retardos usados', '# de observaciones usadas'])
for key,value in df3_test[4].items():
    salida_df3['Critical Value (%s)'%key] = value
print(salida_df3)

```

Resultados del Test de Dickey Fuller para 1 diferenciación de la serie

Consumo Diesel

Estadístico de prueba	-6.884266e+00
p-value	1.407679e-09
# de retardos usados	1.100000e+01
# de observaciones usadas	2.680000e+02
Critical Value (1%)	-3.454988e+00
Critical Value (5%)	-2.872386e+00
Critical Value (10%)	-2.572549e+00

dtype: float64

El p-valor es menor a 0.05 se rechaza la hipótesis nula, Ha: No hay raíces unitarias se acepta

```


[79]: plt.figure(figsize=(12, 6))
plt.title('Consumo Diesel con diferenciación')
plt.plot(df3_diff)

```

```


[79]: [<matplotlib.lines.Line2D at 0x228e04250d0>]

```

```
[80]: plt.figure(figsize=(12, 4))
df3_log_diff = df3_log.diff()
plt.title('Consumo de Diesel con Transformación logarítmica y diferenciación')
plt.plot(df3_log_diff)
```

[80]: [




```
[81]: df3_log_diff.dropna(inplace = True)
df3_log_diff_acf = acf(df3_log_diff,nlags=5,fft=False)
df3_log_diff_pacf = pacf(df3_log_diff, nlags=36)
df3_log_diff_acf, df3_log_diff_pacf
```

```
[81]: (array([ 1.          , -0.19722102,  0.02907365,  0.01670783, -0.06106934,
            -0.0078728 ]),
      array([ 1.          , -0.19792791, -0.01029599,  0.02155302, -0.05637933,
            -0.03308451, -0.46748652, -0.38169926, -0.21806094, -0.07405124,
            -0.11861396, -0.21578912,  0.23083654,  0.04752106,  0.03563811,
            0.01913844, -0.02180888, -0.02791521, -0.20400601, -0.24574491,
            -0.14732467, -0.13522994, -0.21210924, -0.03270305,  0.15932065,
            -0.06794954,  0.05472767,  0.01121561, -0.06728093,  0.07590124,
            -0.25579207, -0.12546835,  0.06216934, -0.09476825,  0.05986017,
            -0.14906471,  0.13643478]))
```

```
[82]: plt.figure(figsize=(10, 6))
      plot_acf(df3_log_diff, lags=36)
      plt.xlabel('Lags')
      plt.ylabel('Autocorrelación')
      plt.title('Análisis ACF "Consumo Diesel"')
      plt.show()
```

<Figure size 1000x600 with 0 Axes>



Se sugiere un valor de desde de 2 para el Q Parece haber un componente de temporadas anual por las correlaciones de 12 24 y 36 meses


```
[83]: plt.figure(figsize=(8, 6))
      plot_pacf(df3_log_diff, lags=12)
      plt.xlabel('Lags')
      plt.ylabel('Autocorrelación Parcial')
      plt.title('Análisis PACF "Consumo Diesel"')
      plt.show()
```

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\graphics\tsaplots.py:348:
FutureWarning: The default method 'yw' can produce PACF values outside of the

$[-1,1]$ interval. After 0.13, the default will change to unadjusted Yule-Walker ('ywm'). You can use this method now by setting method='ywm'.

```
warnings.warn(
```

<Figure size 800x600 with 0 Axes>



para el P podemos usar 2 también

```
[84]: df3_log_D = df3_log.diff(12)
      df3_log_D.dropna(inplace=True)
```


```
[85]: plt.figure(figsize=(8, 6))
      plot_pacf(df3_log_D, lags=6)
      plt.xlabel('Lags')
      plt.ylabel('Autocorrelación Parcial')
      plt.title('Análisis PACF "Consumo Diesel"')
      plt.show()
```

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\graphics\tsaplots.py:348:

FutureWarning: The default method 'yw' can produce PACF values outside of the $[-1,1]$ interval. After 0.13, the default will change to unadjusted Yule-Walker ('ywm'). You can use this method now by setting method='ywm'.

```
warnings.warn(
```

<Figure size 800x600 with 0 Axes>



Luego de hacer una diferenciación estacional, podemos observar que prácticamente se anulan los coeficientes después de $p=3$. Probemos el siguiente componente estacional: - $P = 3$ - $D = 1$ - $Q = 0$

```
[86]: df3_modelo212 = SARIMAX(df3_log, order=(2,1,2), seasonal_order=(3,1,2,12),
    ↪ enforce_stationarity=False, enforce_invertibility=False)
df3_resultado_m212 = df3_modelo212.fit()
print(df3_resultado_m212.summary().tables[1])
```

```
C:\ProgramData\Anaconda3\lib\site-
packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency
information was provided, so inferred frequency MS will be used.
    self._init_dates(dates, freq)
C:\ProgramData\Anaconda3\lib\site-
packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency
information was provided, so inferred frequency MS will be used.
    self._init_dates(dates, freq)
C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\base\model.py:604:
ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check
mle_retvals
    warnings.warn("Maximum Likelihood optimization failed to "
```

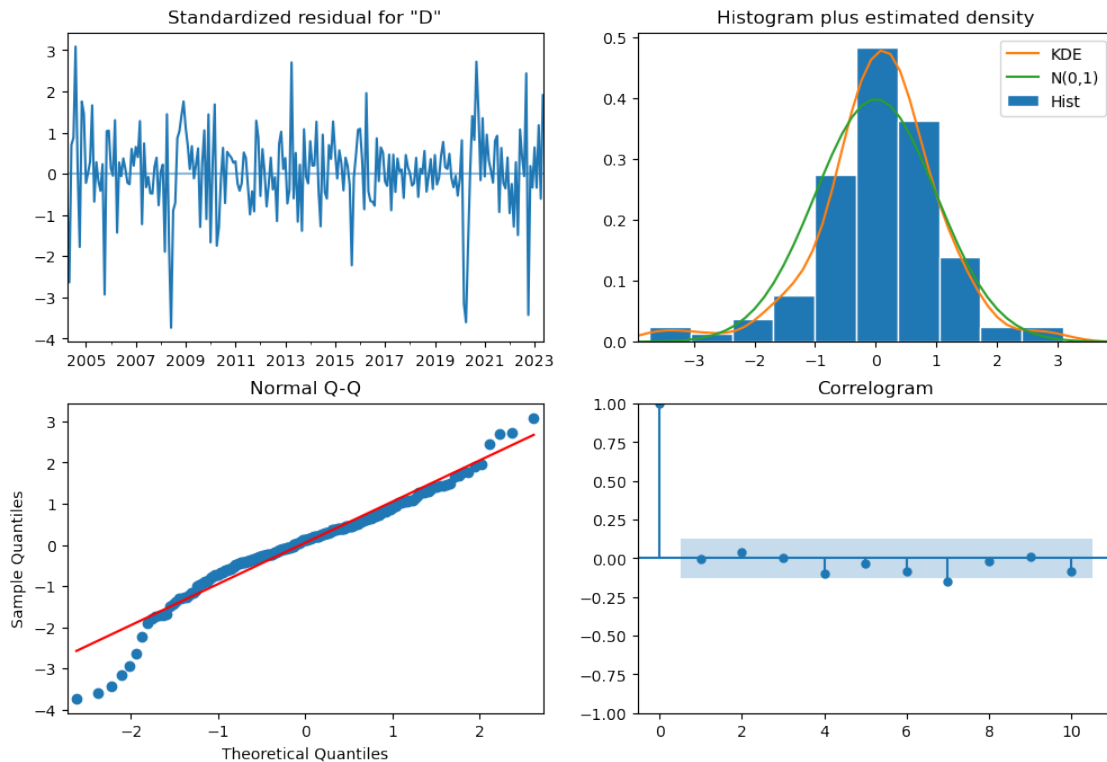
	coef	std err	z	P> z	[0.025	0.975]
ar.L1	-0.6773	0.263	-2.572	0.010	-1.194	-0.161
ar.L2	-0.0417	0.128	-0.325	0.745	-0.293	0.210
ma.L1	1.0071	0.411	2.451	0.014	0.202	1.812
ma.L2	-0.6462	0.303	-2.131	0.033	-1.241	-0.052
ar.S.L12	-0.2644	0.203	-1.303	0.192	-0.662	0.133
ar.S.L24	-0.1147	0.102	-1.121	0.262	-0.315	0.086
ar.S.L36	-0.0441	0.082	-0.540	0.589	-0.204	0.116
ma.S.L12	-0.7713	2.443	-0.316	0.752	-5.559	4.017

ma.S.L24	-0.2246	0.625	-0.359	0.719	-1.449	1.000
sigma2	0.0018	0.004	0.436	0.663	-0.006	0.010

=====

Análisis de residuos

```
[87]: df3_resultado_m212.plot_diagnostics(figsize=(12, 8))
plt.show()
```



Como se puede observar los residuos presentan una distribución masomenos normal, y parecen comportarse como un ruido blanco dado que en el correlograma no hay autocorrelaciones significativas

```
[88]: print("Resultados de AIC (Akaike information criterion)")
print("Modelo 121=",df3_resultado_m212.aic)
print("Modelo 221=",df3_resultado_m212.aic)
```

```
Resultados de AIC (Akaike information criterion)
Modelo 121= -577.3101094727722
Modelo 221= -577.3101094727722
```

Predicción con el mejor modelo

```
[89]: # Generate forecast for the next 12 months
forecast_steps = 12
```

```


forecast = df3_resultado_m212.forecast(steps=forecast_steps)
forecast_original_scale = np.exp(forecast)

# Create a date range for the forecasted months
last_date = df3.index[-1] # Assuming 'df' is your original data DataFrame
forecast_index = pd.date_range(start=last_date, periods=forecast_steps + 1,
    ↪freq='M')[1:]

# Create a DataFrame to hold the forecasted values
forecast_df = pd.DataFrame({'Forecast': forecast}, index=forecast_index)

# Plot the original data and the forecasted values
plt.figure(figsize=(10, 6))
plt.plot(df3.index, df3.values, label='Original Data')
plt.plot(forecast_df.index, forecast_original_scale, label='Forecast',
    ↪color='orange')
plt.xlabel('Date')
plt.ylabel('Value')
plt.title('SARIMAX para 12 meses Diesel')
plt.legend()
plt.show()

```



1.1 Comparacion con LSTM

1.1.1 Importacion de gasolina regular

```
[90]: # Split data into train and test sets
train_size = int(0.8 * len(df1))
train_df = df1[:train_size]
test_df = df1[train_size:]

# ARIMA model
order = (5, 1, 3) # Example ARIMA order
arima_model = ARIMA(train_df, order=order)
arima_results = arima_model.fit()

# LSTM model
def create_lstm_model():
    model = Sequential()
    model.add(LSTM(50, input_shape=(1, 1)))
    model.add(Dense(15, activation='relu'))
    model.add(Dense(1))
    model.compile(optimizer='adam', loss='mean_squared_error')
    return model

X_train = np.array(train_df).reshape(-1, 1, 1)
y_train = np.array(train_df)

lstm_model = create_lstm_model()
lstm_model.fit(X_train, y_train, epochs=50, batch_size=32, verbose=0)

# ARIMA
arima_predictions = arima_results.forecast(steps=len(test_df))

# LSTM
lstm_predictions = []
for i in range(len(test_df)):
    lstm_input = np.array([lstm_predictions[-1]] if lstm_predictions else
        ↪ X_train[-1])
    lstm_pred = lstm_model.predict(lstm_input.reshape(1, 1, 1))
    lstm_predictions.append(lstm_pred[0, 0])

# Calculate RMSE and MAPE for both models
rmse_arima = np.sqrt(mean_squared_error(test_df, arima_predictions))
mape_arima = np.mean(np.abs((test_df - arima_predictions) / test_df)) * 100

rmse_lstm = np.sqrt(mean_squared_error(test_df, lstm_predictions))
mape_lstm = np.mean(np.abs((test_df - lstm_predictions) / test_df)) * 100

# Print the results
```

```

print("ARIMA - RMSE:", rmse_arima)
print("ARIMA - MAPE:", mape_arima)
print("LSTM - RMSE:", rmse_lstm)
print("LSTM - MAPE:", mape_lstm)
print("STD:", df1.std())

```

```

C:\ProgramData\Anaconda3\lib\site-
packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency
information was provided, so inferred frequency MS will be used.
    self._init_dates(dates, freq)
C:\ProgramData\Anaconda3\lib\site-
packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency
information was provided, so inferred frequency MS will be used.
    self._init_dates(dates, freq)
C:\ProgramData\Anaconda3\lib\site-
packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency
information was provided, so inferred frequency MS will be used.
    self._init_dates(dates, freq)
C:\ProgramData\Anaconda3\lib\site-
packages\statsmodels\tsa\statespace\sarimax.py:966: UserWarning: Non-stationary
starting autoregressive parameters found. Using zeros as starting parameters.
    warn('Non-stationary starting autoregressive parameters')
C:\ProgramData\Anaconda3\lib\site-
packages\statsmodels\tsa\statespace\sarimax.py:978: UserWarning: Non-invertible
starting MA parameters found. Using zeros as starting parameters.
    warn('Non-invertible starting MA parameters found.')

1/1 [=====] - 1s 660ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 33ms/step

```



```

1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 57ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 78ms/step
1/1 [=====] - 0s 37ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 38ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 55ms/step
1/1 [=====] - 0s 46ms/step
1/1 [=====] - 0s 45ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 51ms/step
1/1 [=====] - 0s 40ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 26ms/step

```

ARIMA - RMSE: 200348.75221223445

ARIMA - MAPE: 23.55410108988794

LSTM - RMSE: 728430.4762709666

LSTM - MAPE: 99.999183747892

STD: 219364.3548742019

Para esta primera serie, vemos que el rmse es significativamente mejor en el arima que el LSTM utilizado. De hecho este arima es aceptable considerando que un MAPE de 23.55% no esta nada mal considerando el valor de la varianza, ademas de que lo hicimos con el test. Por lo que utilizaremos ARIMA para predecir la

1.1.2 Precios de gasolina superior

```
[91]: # Split data into train and test sets
train_size = int(0.8 * len(df2))
train_df = df2[:train_size]
test_df = df2[train_size:]

# ARIMA model
order = (0, 1, 0) # Example ARIMA order
arima_model = ARIMA(train_df, order=order)
arima_results = arima_model.fit()

# LSTM model
def create_lstm_model():
    model = Sequential()
    model.add(LSTM(50, input_shape=(1, 1)))
    model.add(Dense(15, activation='relu'))
    model.add(Dense(1))
    model.compile(optimizer='adam', loss='mean_squared_error')
    return model

X_train = np.array(train_df).reshape(-1, 1, 1)
y_train = np.array(train_df)

lstm_model = create_lstm_model()
lstm_model.fit(X_train, y_train, epochs=50, batch_size=32, verbose=0)

# ARIMA
arima_predictions = arima_results.forecast(steps=len(test_df))

# LSTM
lstm_predictions = []
for i in range(len(test_df)):
    lstm_input = np.array([lstm_predictions[-1]] if lstm_predictions else
        ↪ X_train[-1])
    lstm_pred = lstm_model.predict(lstm_input.reshape(1, 1, 1))
    lstm_predictions.append(lstm_pred[0, 0])

# Calculate RMSE and MAPE for both models
rmse_arima = np.sqrt(mean_squared_error(test_df, arima_predictions))
mape_arima = np.mean(np.abs((test_df - arima_predictions) / test_df)) * 100

rmse_lstm = np.sqrt(mean_squared_error(test_df, lstm_predictions))
mape_lstm = np.mean(np.abs((test_df - lstm_predictions) / test_df)) * 100

# Print the results
print("ARIMA - RMSE:", rmse_arima)
```

```

print("ARIMA - MAPE:", mape_arima)
print("LSTM - RMSE:", rmse_lstm)
print("LSTM - MAPE:", mape_lstm)
print("STD", df2.std())

```

```

C:\ProgramData\Anaconda3\lib\site-
packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency
information was provided, so inferred frequency D will be used.

```

```

    self._init_dates(dates, freq)

```

```

C:\ProgramData\Anaconda3\lib\site-
packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency
information was provided, so inferred frequency D will be used.

```

```

    self._init_dates(dates, freq)

```

```

C:\ProgramData\Anaconda3\lib\site-
packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency
information was provided, so inferred frequency D will be used.

```

```

    self._init_dates(dates, freq)

```

```

1/1 [=====] - 1s 582ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 83ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 29ms/step

```

```

1/1 [=====] - 0s 51ms/step
1/1 [=====] - 0s 49ms/step
1/1 [=====] - 0s 61ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 59ms/step
1/1 [=====] - 0s 46ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 47ms/step
1/1 [=====] - 0s 40ms/step
1/1 [=====] - 0s 38ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 29ms/step

```

```

1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 44ms/step
1/1 [=====] - 0s 38ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 38ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 40ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 39ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 37ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 38ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 39ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 24ms/step

```

1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 38ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 37ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 37ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 28ms/step

```

1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 31ms/step
ARIMA - RMSE: 1.2877411419525537
ARIMA - MAPE: 3.340024289795889
LSTM - RMSE: 2.4154358967476734
LSTM - MAPE: 6.602520454230637
STD 4.555857482383936

```

En este caso el ARIMA tiene un muy buen rendimiento en el test ya que el rmse casi 4 veces mejor que la desviacion estandar, mientras que el LSTM es simplemente aceptable. Por lo tanto nos quedaremos con el ARIMA.

1.1.3 Consumos de diesel

```

[92]: # Split data into train and test sets
train_size = int(0.8 * len(df3))
train_df = df3[:train_size]
test_df = df3[train_size:]

# ARIMA model
order = (2,1,2) # Example ARIMA order
arima_model = ARIMA(train_df, order=order)
arima_results = arima_model.fit()

# LSTM model
def create_lstm_model():
    model = Sequential()
    model.add(LSTM(50, input_shape=(1, 1)))
    model.add(Dense(15, activation='relu'))
    model.add(Dense(1))
    model.compile(optimizer='adam', loss='mean_squared_error')
    return model

X_train = np.array(train_df).reshape(-1, 1, 1)
y_train = np.array(train_df)

```

```

lstm_model = create_lstm_model()
lstm_model.fit(X_train, y_train, epochs=50, batch_size=32, verbose=0)

# ARIMA
arima_predictions = arima_results.forecast(steps=len(test_df))

# LSTM
lstm_predictions = []
for i in range(len(test_df)):
    lstm_input = np.array([lstm_predictions[-1]] if lstm_predictions else
↳X_train[-1])
    lstm_pred = lstm_model.predict(lstm_input.reshape(1, 1, 1))
    lstm_predictions.append(lstm_pred[0, 0])

# Calculate RMSE and MAPE for both models
rmse_arima = np.sqrt(mean_squared_error(test_df, arima_predictions))
mape_arima = np.mean(np.abs((test_df - arima_predictions) / test_df)) * 100

rmse_lstm = np.sqrt(mean_squared_error(test_df, lstm_predictions))
mape_lstm = np.mean(np.abs((test_df - lstm_predictions) / test_df)) * 100

# Print the results
print("ARIMA - RMSE:", rmse_arima)
print("ARIMA - MAPE:", mape_arima)
print("LSTM - RMSE:", rmse_lstm)
print("LSTM - MAPE:", mape_lstm)
print("STD", df3.std())

```

```

C:\ProgramData\Anaconda3\lib\site-
packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency
information was provided, so inferred frequency MS will be used.
    self._init_dates(dates, freq)
C:\ProgramData\Anaconda3\lib\site-
packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency
information was provided, so inferred frequency MS will be used.
    self._init_dates(dates, freq)
C:\ProgramData\Anaconda3\lib\site-
packages\statsmodels\tsa\base\tsa_model.py:471: ValueWarning: No frequency
information was provided, so inferred frequency MS will be used.
    self._init_dates(dates, freq)
C:\ProgramData\Anaconda3\lib\site-
packages\statsmodels\tsa\statespace\sarimax.py:966: UserWarning: Non-stationary
starting autoregressive parameters found. Using zeros as starting parameters.
    warn('Non-stationary starting autoregressive parameters')
C:\ProgramData\Anaconda3\lib\site-
packages\statsmodels\tsa\statespace\sarimax.py:978: UserWarning: Non-invertible

```


starting MA parameters found. Using zeros as starting parameters.
warn('Non-invertible starting MA parameters found.')

```
1/1 [=====] - 1s 836ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 37ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 38ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 38ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 26ms/step
```

```

1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 42ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 28ms/step
ARIMA - RMSE: 195045.63885886592
ARIMA - MAPE: 14.201927554073238
LSTM - RMSE: 1146160.8889994498
LSTM - MAPE: 99.99963639993777
STD 205345.61556490546

```

En este caso el ARIMA sigue siendo superior al LSTM. Aún así no es muy bueno ninguno de ambos modelos. De elegir nos quedaríamos con el ARIMA.

1.2 Discusión inciso 6

Creemos que las predicciones sí son afectas por la pandemia en particular. Esto se puede apreciar con mayor detalle en la predicción del consumo de Diesel ya que la caída pronunciada en este provoca que los cambios en la proyección del arima sean más pronunciados intentando reflejar la tendencia de los datos históricos. De no ser así la variación se debería observar menos pronunciada según la época del año.